

# Compressed Sensing

Practical Optimisation

Abhishek Shenoy

## Abstract

This report investigates the use of convex optimisation to solve various norm approximation problems. We analyse various optimisation methods using the central path method after formulating them as linear programming problems.

## 1 Norm Approximation

The simplest norm approximation problem is an unconstrained problem of the following form.

$$\min ||\mathbf{Ax} - \mathbf{b}||$$

$\mathbf{A} \in \mathbb{R}^{m \times n}$  and  $\mathbf{b} \in \mathbb{R}^m$  are given as problem data,  $\mathbf{x} \in \mathbb{R}^n$  is the variable and  $||\cdot||$  is a norm on  $\mathbb{R}^m$ . The dataset used for norm approximations below have dimension  $m = 2n$ .

### 1.1 Norm Approximation Problems

We investigate the norm approximation problem for the  $l_1$ ,  $l_2$  and  $l_\infty$  norm. The norm definitions for a vector  $\mathbf{y} \in \mathbb{R}^m$  are given by the following:

$$\begin{aligned} l_1(\mathbf{y}) &= ||\mathbf{y}||_1 = \sum_{i=1}^m |y_i| \\ l_2(\mathbf{y}) &= ||\mathbf{y}||_2 = \left( \sum_{i=1}^m y_i^2 \right)^{\frac{1}{2}} \\ l_\infty(\mathbf{y}) &= ||\mathbf{y}||_\infty = \max_{i \in \{1, m\}} |y_i| \end{aligned}$$

Both  $L_1$  and  $L_\infty$  norm approximations have no closed-form solutions. In fact  $L_2$  is the only case with an analytic solution.

In the below cases,  $\tilde{\mathbf{A}} \in \mathbb{R}^{2m \times (n+m)}$ ,  $\tilde{\mathbf{x}} \in \mathbb{R}^{n+m}$ ,  $\tilde{\mathbf{c}} \in \mathbb{R}^{n+m}$ ,  $\tilde{\mathbf{b}} \in \mathbb{R}^{2m}$ .

#### 1.1.1 $L_1$ Norm Approximation

$$\begin{aligned} \min_{\mathbf{x}} ||\mathbf{Ax} - \mathbf{b}||_1 &= \min_{\mathbf{x}} \sum_{i=1}^m |\mathbf{a}_i^T \mathbf{x} - b_i| \\ &= \min_{\mathbf{x}, t_i} \sum_{i=1}^m t_i \quad \{-t_i \leq \mathbf{a}_i^T \mathbf{x} - b_i \leq t_i\}_{i=1}^m \end{aligned}$$

This can be solved via the following linear program:

$$\begin{aligned} \min_{\tilde{\mathbf{x}}} \tilde{\mathbf{c}}^T \tilde{\mathbf{x}} \\ \tilde{\mathbf{A}} \tilde{\mathbf{x}} \leq \tilde{\mathbf{b}} \end{aligned}$$

$$\tilde{\mathbf{x}} = \begin{bmatrix} \mathbf{x} \\ \mathbf{t} \end{bmatrix} \quad \tilde{\mathbf{c}} = \begin{bmatrix} \mathbf{0} \\ \mathbf{1} \end{bmatrix} \quad \tilde{\mathbf{A}} = \begin{bmatrix} \mathbf{A} & -\mathbf{I} \\ -\mathbf{A} & -\mathbf{I} \end{bmatrix} \quad \tilde{\mathbf{b}} = \begin{bmatrix} \mathbf{b} \\ -\mathbf{b} \end{bmatrix}$$

#### 1.1.2 $L_\infty$ Norm Approximation

$$\begin{aligned} \min_{\mathbf{x}} ||\mathbf{Ax} - \mathbf{b}||_\infty &= \min_{\mathbf{x}} \max_i |\mathbf{a}_i^T \mathbf{x} - b_i| \\ &= \min_{\mathbf{x}, t} t \quad \{-t \leq \mathbf{a}_i^T \mathbf{x} - b_i \leq t\}_{i=1}^m \end{aligned}$$

This can be solved via the following linear program:

$$\begin{aligned} \min_{\tilde{\mathbf{x}}} \tilde{\mathbf{c}}^T \tilde{\mathbf{x}} \\ \tilde{\mathbf{A}} \tilde{\mathbf{x}} \leq \tilde{\mathbf{b}} \end{aligned}$$

$$\tilde{\mathbf{x}} = \begin{bmatrix} \mathbf{x} \\ t \end{bmatrix} \quad \tilde{\mathbf{c}} = \begin{bmatrix} \mathbf{0} \\ 1 \end{bmatrix} \quad \tilde{\mathbf{A}} = \begin{bmatrix} \mathbf{A} & -\mathbf{1} \\ -\mathbf{A} & -1 \end{bmatrix} \quad \tilde{\mathbf{b}} = \begin{bmatrix} \mathbf{b} \\ -\mathbf{b} \end{bmatrix}$$

#### 1.1.3 $L_2$ Norm Approximation

The  $L_2$  norm approximation problem can be formulated using matrices and vectors leading to an analytic solution for the convex quadratic function.

$$\begin{aligned} \min_{\mathbf{x}} ||\mathbf{Ax} - \mathbf{b}||_2 &= \min_{\mathbf{x}} ||\mathbf{Ax} - \mathbf{b}||_2^2 \\ &= \min_{\mathbf{x}} (\mathbf{Ax} - \mathbf{b})^T (\mathbf{Ax} - \mathbf{b}) \end{aligned}$$

By differentiating with respect to  $\mathbf{x}$ , we can find the least-squares solution.

$$\begin{aligned} \frac{d}{d\mathbf{x}} (\mathbf{Ax} - \mathbf{b})^T (\mathbf{Ax} - \mathbf{b}) &= 0 \\ \frac{d}{d\mathbf{x}} (\mathbf{x}^T \mathbf{A}^T \mathbf{Ax} - 2\mathbf{x}^T \mathbf{A}^T \mathbf{b} + \mathbf{b}^T \mathbf{b}) &= 0 \\ 2(\mathbf{A}^T \mathbf{Ax} - \mathbf{A}^T \mathbf{b}) &= 0 \\ \therefore \mathbf{A}^T \mathbf{Ax} &= \mathbf{A}^T \mathbf{b} \\ \mathbf{x} &= (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b} \end{aligned}$$

## 1.2 Norm Comparisons

A, b Dataset	Size n	Minimised Values			Runtime (ms)		
		$l_1$	$l_2$	$l_\infty$	$l_1$	$l_2$	$l_\infty$
(A <sub>1</sub> , b <sub>1</sub> )	16	12.48	2.1	0.77	8.1	0.9	10.0
(A <sub>2</sub> , b <sub>2</sub> )	64	52.45	5.45	0.82	52.1	1.4	49.5
(A <sub>3</sub> , b <sub>3</sub> )	256	194.88	9.36	0.80	881.5	41.0	775.2
(A <sub>4</sub> , b <sub>4</sub> )	512	402.03	13.86	0.80	8288.3	129.0	6389.1
(A <sub>5</sub> , b <sub>5</sub> )	1024	808.43	18.71	0.79	56359.6	493.0	44112.0

Table 1. Minimisation of different norms of  $||\mathbf{Ax} - \mathbf{b}||$

The minimised values are of the correct magnitude as we would expect the  $l_1$  norm to be the largest as it is the absolute sum,  $l_2$  norm to be the next largest as it is the square-root of the sum of the squares and  $l_\infty$  norm to be the smallest as it is the maximum component of a vector.

The least-squares solution is by far the easiest and quickest to compute due to the analytical solution rather than taking an optimisation approach.

$l_1$  has a relatively faster runtime as compared to  $l_\infty$ . It is also possible to see the non-linearity in the increase in runtime as  $n$  increases. For  $n = 256, 512, 1024$ , we see that for every factor 2 increase in  $n$  there is an approximate factor 4 increase in runtime for  $l_2$  however for  $l_1$  and  $l_\infty$  the increase in runtime is approximately by factor 8. We do however see a rough factor of  $16(4^2)$  increase in runtime for all the 3 norms between  $n = 64$  to  $n = 256$  which is as expected. Below  $n = 64$  all the norms can be computed relatively fast such that the speed difference is of the expected factor but the runtimes are so small that the differences are wide.  $l_2$  is able to compute extremely fast even for  $n = 64$  almost at the same speed as that of  $n = 16$  possibly taking advantage of precompiled Numpy libraries for least-squares.

### 1.3 Norm Residuals

We plot the histogram of the residuals of the norm approximation problem for the data pair  $(\mathbf{A}_5, \mathbf{b}_5)[n = 1024]$  for the  $l_1, l_2, l_\infty$ -norms.

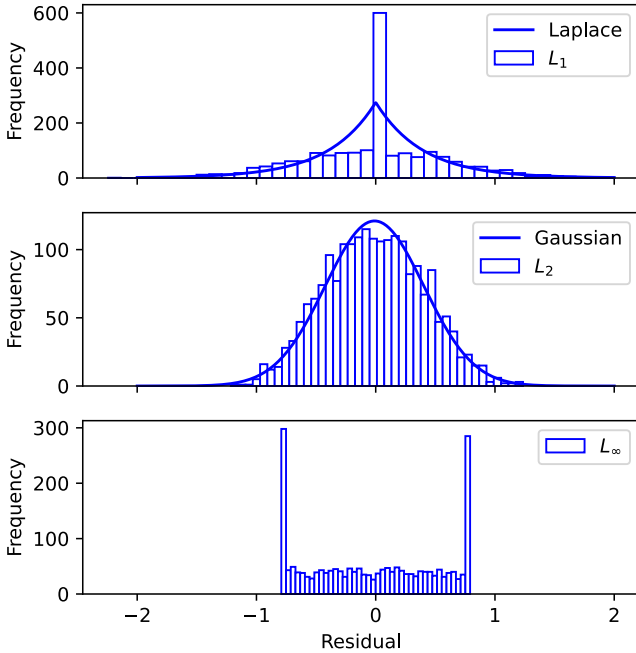


Figure 1. Histogram of the residuals of the norm approximation problem for  $(\mathbf{A}_5, \mathbf{b}_5)$

The amplitude distribution of the optimal residual for the  $l_1$ -norm approximation problem will tend to have more zero and very small residuals, compared to the  $l_2$ -norm approximation solution. In contrast, the  $l_2$ -norm solution will tend to have relatively fewer large residuals (since large residuals incur a

much larger penalty in  $l_2$ -norm approximation than in  $l_1$ -norm approximation). This is visibly clear in Figure 1 as the Laplace distribution for  $l_1$  has a much higher density at 0 than the Gaussian distribution and the fat tails of the Laplace also means that there is a greater number of large residuals for  $l_1$  than for  $l_2$  due to the penalisation by the  $l_2$ -norm. The  $l_2$ -norm approximation has many modest residuals, and relatively few larger ones.

In the case of the  $l_\infty$ -norm, we get a uniform distribution with heavy weightings on the largest outliers. This is explicable as the  $l_\infty$ -norm returns the maximum component.

If you minimise the negative log-likelihood of the error distribution, you get the best estimator. For a Gaussian distribution, this results in a sum of squares minimization. For Laplace, this results in a sum of absolute value minimization. Hence a good fit for the data for a given norm occurs when the residual distribution matches the distribution corresponding to the norm used.

## 2 Central Path Formulation

Consider the following optimisation problem:

$$\begin{aligned} \min_{\mathbf{x}} f_0(\mathbf{x}) \\ \text{subject to } f_i(\mathbf{x}) \leq 0 \quad i = 1, \dots, m \end{aligned}$$

$f_0, f_i$  are convex and twice continuously differentiable functions for  $i = 1, \dots, m$ .

The central path is the solution of the unconstrained minimisation of the following problem:

$$\min_{\mathbf{x}} t f_0(\mathbf{x}) + \phi(\mathbf{x}) \quad t \geq 0$$

$\phi$  is the logarithmic barrier function of the feasibility set  $S = \{\mathbf{x} : f_i(\mathbf{x}) \leq 0\}$ .

### 2.1 $L_1$ CP Formulation

We consider the central path LP formulation of the  $L_1$  norm approximation problem expressed in Section 1.1.1 using the matrix  $\tilde{\mathbf{A}}$  and vectors  $\tilde{\mathbf{x}}, \tilde{\mathbf{b}}, \tilde{\mathbf{c}}$  previously defined.

$$\tilde{\mathbf{x}}^*(t) = \min_{\tilde{\mathbf{x}}} t(\tilde{\mathbf{c}}^T \tilde{\mathbf{x}}) - \underbrace{\sum_{i=1}^m \ln(\tilde{b}_i - \tilde{\mathbf{a}}_i^T \tilde{\mathbf{x}})}_{+\phi(\tilde{\mathbf{x}})}$$

A unique  $\tilde{\mathbf{x}}^*(t)$  exists for all  $t > 0$ . This set of values form the central path from the centre of the polyhedron at  $t = 0$  to the edges of the polyhedron at  $t = \infty$ . The barrier function corresponding to a constraint  $\tilde{\mathbf{A}}\tilde{\mathbf{x}} \leq \tilde{\mathbf{b}}$  forms the boundary of the polyhedron.

The gradient of the cost for a fixed  $t \geq 0$  can be calculated by differentiating with respect to  $\tilde{\mathbf{x}}$  as the following.

$$\text{Gradient} \quad t\tilde{\mathbf{c}} + \underbrace{\sum_{i=1}^m \frac{1}{\tilde{b}_i - \tilde{\mathbf{a}}_i^T \tilde{\mathbf{x}}} \tilde{\mathbf{a}}_i}_{\nabla \phi(\tilde{\mathbf{x}})}$$

Defining a vector  $\mathbf{d}$  for which  $d_i = \frac{1}{b_i - \tilde{\mathbf{a}}_i^T \tilde{\mathbf{x}}}$ , the gradient can be simplified to the below.

$$\text{Gradient } t\tilde{\mathbf{c}} + \tilde{\mathbf{A}}^T \mathbf{d}$$

The solution for the non- $t$  term  $\tilde{\mathbf{A}}^T \mathbf{d} = 0$  where  $\mathbf{d}$  is a function of  $\mathbf{x}$  gives the analytic centre  $\mathbf{x}_{ac}$ . The term with  $t$  essentially creates a path with the variation of  $t$  known as the central path.  $\mathbf{x}_{ac}$  exists and is unique if and only if the polyhedron (Dikin ellipsoid) is non-empty and bounded (Vandenberghe, 2013).

As we move along the central path, the contour gets closer to the actual polyhedron. Because of this geometry, the log-barrier method is called an interior-point method, since we start at an interior point and move along the central path to get to the solution.

## 2.2 Backtracking LineSearch

We can solve for the solution by minimising the gradient (setting the gradient to 0). For this particular case, we consider  $t = 1$  corresponding to the pair  $(\mathbf{A}_3, \mathbf{b}_3)[n = 256]$ . We can use many methods to solve the problem however here we specifically focus on using a first-order gradient method with backtracking linesearch to analyse the convergence properties.

---

### Algorithm 1 Backtracking LineSearch

---

**given** a descent direction  $\Delta \mathbf{x}$  for  $f$  at  $\mathbf{x} \in \text{dom} f$   
 $\alpha \in (0, 0.5), \beta \in (0, 1)$   
 $t \leftarrow 1$   
**while**  $f(\mathbf{x} + t \Delta \mathbf{x}) > f(\mathbf{x}) + \alpha t \nabla f(\mathbf{x})^T \Delta \mathbf{x}$  **do**  
 $t \leftarrow \beta t$   
**end while**

---

The method for the backtracking linesearch has been shown in Algorithm 1 for which there are several parameters that we need to pick carefully (Boyd and Vandenberghe, 2004).

This specific problem is convex and differentiable, and  $\nabla f$  is Lipschitz continuous with constant  $L > 0$ . Therefore gradient descent can be used. Gradient descent with backtracking for a given  $\beta$  satisfies the following inequality:

$$f(\mathbf{x}_k) - \underbrace{f(\mathbf{x}^*)}_{p^*} \leq \frac{\|\mathbf{x}_0 - \mathbf{x}^*\|^2}{2t_{min}k} \quad t_{min} = \min(1, \beta/L)$$

Gradient descent with a fixed step-size  $t$  (ie.  $t_{min} = t$ ) has convergence rate  $O(1/k)$  for  $k$  updates of  $\mathbf{x}$ . However with backtracking for  $t_{min} = \min(1, \beta/L)$ , this is only true if  $\beta$  is not too small so that we do not lose much compared to fixed step size (Gordon and Tibshirani, 2012).

### 2.2.1 Hyperparameters

$\alpha$  is typically chosen between 0.01 and 0.5, meaning that we accept a decrease in  $f$  between 1% and 50% of the prediction based on the linear extrapolation.  $\beta$  is often chosen to be between 0.1 (corresponding to a very crude search) and 0.8 (corresponding to a less crude search).




Figure 2. Linear convergence of backtracking linesearch on semi-log plot

## 2.3 Convergence Analysis

Figure 2 shows linearity up to the point of convergence using the gradient method with backtracking linesearch where  $p^*$  is the local optimum reached when the initialisation is  $t = 1$ .

As the number of iterations increases  $k \rightarrow \infty$ , the error must tend toward 0 as  $|f(\mathbf{x}_k) - p^*| \rightarrow 0$  and hence  $f(\mathbf{x}_k) - p^* \leq d \cdot c^k$  for some constants  $c$  and  $d$ . More specifically this is given by the following:

$$\frac{f(\mathbf{x}_k) - p^*}{f(\mathbf{x}_0) - p^*} \leq c^k$$

Therefore the error  $f(\mathbf{x}_k) - p^*$  converges to zero at least as fast as a geometric series. In the context of iterative numerical methods, this is called linear convergence since it is linear on a log-linear plot of error versus iteration number.

For the convex function  $f$  where  $m\mathbf{I} \preceq \nabla^2 f \preceq M\mathbf{I}$  for the positive constants  $m$  and  $M$ , the bound for the condition number of the Hessian is given by  $M/m$ . For the case of gradient descent with exact linesearch,  $c = 1 - m/M$ . For gradient descent with backtracking linesearch,  $c = 1 - \min(2m\alpha, 2\beta\alpha m/M)$ . Therefore it is clear that if the condition number of the Hessian is too large, then the rate of convergence slows down. Nonetheless the convergence is at least linear (in the sense of iterative methods).

Note however that backtracking does not effectively find the global solution for the formulated problem as in Section 1.1.1 as  $t$  varies from 0 to  $\infty$ . By checking the condition number of the Hessian  $\nabla^2 \phi(\tilde{\mathbf{x}}) = \mathbf{A}^T \text{diag}(\mathbf{d})^2 \mathbf{A}$ , we find that it is extremely large hence explaining the slow convergence (Boyd and Vandenberghe (2004) Pg 475).

## 3 LASSO Problem

Consider the following  $l_1$ -regularised least squares problem where  $\lambda > 0$  is the regularisation parameter and  $t$  varies from 0 to  $\infty$ :

$$\min_{\mathbf{x}} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|_2^2 + \lambda \|\mathbf{x}\|_1 \quad (1)$$

This problem is also known as the basis pursuit denoising problem (BPDN) / least absolute shrinkage and selection operator (LASSO) which always has a solution but is not necessarily unique.

Since there is no analytical solution, we need to express this either as a differentiable optimisation problem or using an iterative method to compute the solution.

### 3.1 Convex Quadratic Formulation

Equation 1 can be transformed to a convex quadratic problem with linear inequality constraints as shown in Equation 2.

By considering a new variable  $u_i = |x_i|$  we define the inequality constraint as  $-x_i < u_i < x_i$  which we can now use to define the logarithmic barrier function. Since there are two inequalities, there is a requirement of two barrier functions which we can combine into a single barrier function.

$$\begin{aligned}\phi_t(\mathbf{x}, \mathbf{u}) &= t \left( \|\mathbf{Ax} - \mathbf{b}\|_2^2 + \lambda \sum_{i=1}^n u_i \right) + \Phi(\mathbf{x}, \mathbf{u}) \\ &= t \left( \|\mathbf{Ax} - \mathbf{b}\|_2^2 + \lambda \mathbf{u}^T \mathbf{1} \right) + \Phi(\mathbf{x}, \mathbf{u}) \\ \Phi(\mathbf{x}, \mathbf{u}) &= - \sum_{i=1}^n \ln(u_i + x_i) - \sum_{i=1}^n \ln(u_i - x_i) \\ &= - \sum_{i=1}^n \ln(u_i^2 - x_i^2)\end{aligned}\quad (2)$$

Note that the equivalent mathematical notation for element-wise operations for  $\Phi(\mathbf{x}, \mathbf{u})$  that can be converted into Python code is given by the following:

$$\Phi(\mathbf{x}, \mathbf{u}) = - \sum_{i=1}^n (\ln \circ (\mathbf{u}^{\circ 2} - \mathbf{x}^{\circ 2}))_i$$

### 3.2 Gradient and Hessian

We can now derive the gradient  $\mathbf{g} = \nabla \phi_t$  and the Hessian  $\mathbf{H} = \nabla^2 \phi_t$  similarly using Hadamard notation for element-wise operations where  $\circ$  denotes element-wise function application,  $\odot$  denotes element-wise multiplication and  $\oslash$  denotes element-wise division.

$$\begin{aligned}\mathbf{g} = \nabla \phi_t(\mathbf{x}, \mathbf{u}) &= \begin{bmatrix} \nabla_{\mathbf{x}} \phi_t(\mathbf{x}, \mathbf{u}) \\ \nabla_{\mathbf{u}} \phi_t(\mathbf{x}, \mathbf{u}) \end{bmatrix} \\ &= \begin{bmatrix} 2t\mathbf{A}^T(\mathbf{Ax} - \mathbf{b}) + 2\mathbf{x} \oslash (\mathbf{u}^{\circ 2} - \mathbf{x}^{\circ 2}) \\ t\lambda \mathbf{1} - 2\mathbf{u} \oslash (\mathbf{u}^{\circ 2} - \mathbf{x}^{\circ 2}) \end{bmatrix}\end{aligned}$$

$$\begin{aligned}\mathbf{H} = \nabla^2 \phi_t(\mathbf{x}, \mathbf{u}) &= t\nabla^2 \|\mathbf{Ax} - \mathbf{b}\|_2^2 + \nabla^2 \Phi(\mathbf{x}, \mathbf{u}) \\ &= t\nabla^2 (\mathbf{Ax} - \mathbf{b})^T (\mathbf{Ax} - \mathbf{b}) + \nabla^2 \Phi(\mathbf{x}, \mathbf{u}) \\ &= \begin{bmatrix} 2t\mathbf{A}^T \mathbf{A} + \mathbf{D}_1 & \mathbf{D}_2 \\ \mathbf{D}_2 & \mathbf{D}_1 \end{bmatrix}\end{aligned}$$

$$\begin{aligned}\mathbf{D}_1 &= \mathbf{diag}(2(\mathbf{u}^{\circ 2} + \mathbf{x}^{\circ 2}) \oslash (\mathbf{u}^{\circ 2} - \mathbf{x}^{\circ 2})^{\circ 2}) \\ &= \mathbf{diag}\left(\frac{2(u_1^2 + x_1^2)}{(u_1^2 - x_1^2)^2}, \dots, \frac{2(u_n^2 + x_n^2)}{(u_n^2 - x_n^2)^2}\right) \\ \mathbf{D}_2 &= \mathbf{diag}\left(-4(\mathbf{u} \odot \mathbf{x}) \oslash (\mathbf{u}^{\circ 2} - \mathbf{x}^{\circ 2})^{\circ 2}\right) \\ &= \mathbf{diag}\left(\frac{-4u_1x_1}{(u_1^2 - x_1^2)^2}, \dots, \frac{-4u_nx_n}{(u_n^2 - x_n^2)^2}\right)\end{aligned}$$

### 3.3 Sparse Signal Reconstruction

We consider a sparse signal recovery problem with a signal  $\mathbf{x}_0 \in \mathbb{R}^{256}$  which consists of 10 spikes of amplitude  $\pm 1$ . The measurement matrix  $\mathbf{A} \in \mathbb{R}^{60 \times 256}$  and  $\mathbf{x}_0$  and the vector of observations is  $\mathbf{b} = \mathbf{Ax}_0$ .

We perform a sparse signal reconstruction by applying a Newton interior-point method to Equation 2 with regularisation parameter  $\lambda = 0.01\lambda_{max}$  with  $\lambda_{max} = |2\mathbf{A}^T \mathbf{b}|_{\infty}$  and plot the original and reconstructed signals (Fig 3).

This problem is an example of **compressed sensing** and solving this is extremely useful in signal processing for efficiently acquiring and reconstructing a signal by finding solutions to under-determined linear systems.

Note that we are explicitly investigating  $\lambda = 0.01\lambda_{max}$  but if we wanted to determine the optimum  $\lambda$ , we could use the Morozov discrepancy principle. Since we have the desired signal, we can determine this optimum value of  $\lambda$  by plotting the variation of the reconstruction error  $\|\mathbf{x} - \mathbf{x}^*\|$  with  $\lambda$  and picking  $\lambda$  which has the minimum error. Of course, in practice this method is not practically viable since we have to know the solution in advance and hence the Morozov discrepancy principle is used (Leykekhman, 2018).

#### 3.3.1 Newton Interior-Point Method

---

**Algorithm 2** Newton IPM for  $l_1$ -regularised LSPs

---

**Given** relative tolerance  $\epsilon > 0, t \leftarrow \frac{1}{\lambda}$   
 $\mathbf{x} = \mathbf{0}, \mathbf{u} = \mathbf{1}$   
**while** Stop Criterion  $\leq \epsilon$  **do**  
    **Compute** search direction  $(\Delta \mathbf{x}, \Delta \mathbf{u})$  using Newton  
    **Compute** step size  $\tau$  by backtracking linesearch  
     $\mathbf{x} \leftarrow \mathbf{x} + \tau \Delta \mathbf{x}$   
     $\mathbf{u} \leftarrow \mathbf{u} + \tau \Delta \mathbf{u}$   
    **Update**  $t$   
**end while**

---

When an iterative method is used to approximately solve the Newton system, the overall method is called a truncated Newton method.

To compute the search direction  $\Delta \mathbf{d}$ , we need to solve the following equation where  $\mathbf{H}$  is the Hessian and  $\mathbf{g}$  is the gradient.

$$\mathbf{H} \underbrace{\begin{bmatrix} \Delta \mathbf{x} \\ \Delta \mathbf{u} \end{bmatrix}}_{\Delta \mathbf{d}} = -\mathbf{g}$$

Instead of using the truncated Newton interior-point method (TNIPM) (approximating the search direction) for solving this task as done in [Kim et al. \(2007\)](#) using preconditioned gradient (PCG) methods ([Huang et al., 2018](#)), we directly compute the search-direction using the Hessian hence this is the full Newton interior-point method defined in Algorithm 2. This is only feasible in this problem because we know that the signal that needs to be reconstructed is relatively small and sparse specifically with only 10 non-zero elements out of the 256 components. Generally it is possible to approximate the Newton gradient and search direction by decomposing it into a matrix for which only the diagonal of the Hessian is required to be computed ([Bickson and Dolev, 2010](#)).

We explored 2 methods for using a NIPM specifically using the barrier method (BM) with Newton as well as the custom interior-point method (CIPM) defined in [Kim et al. \(2007\)](#).

### 3.3.2 $t$ -Update Methods

Updating  $t$  is needed to traverse the central path. This needs to be chosen appropriately although it is very flexible.

We considered two methods of updating  $t$  specifically the use of the update rule  $t = \mu t$  as well as the update rule  $t = \max(\mu \min(\frac{2n}{\eta}, t), t)$  if  $\tau \geq \tau_{min}$  defined in [Kim et al. \(2007\)](#) where  $\eta$  is the duality gap and  $n$  is the dimension of  $\mathbf{x}$ .

#### Barrier Method

The barrier method uses  $t = \mu t$  with  $\mu > 1$  ([Boyd and Vandenberghe \(2004\)](#) Pg 569) as the update rule for  $t$ .

The choice of the parameter  $\mu$  involves a trade-off in the number of inner and outer iterations required. If  $\mu$  is small (near 1) then at each outer iteration  $t$  increases by a small factor and therefore we require fewer Newton steps (inner iterations) to converge to the new centre but require larger number of outer iterations to move to the edge of the polyhedron along the central path. We also have to consider the initial value  $t_0$  for the same reason. The value for  $\mu$  and  $t_0$  can therefore be determined by experimentation ([Boyd and Vandenberghe \(2004\)](#) Pg 570).

#### CIPM

The CIPM that we investigate uses the update rule  $t = \max(\mu \min(\frac{2n}{\eta}, t), t)$  if  $\tau \geq \tau_{min}$  defined in [Kim et al. \(2007\)](#) where  $\eta$  is the duality gap and  $n$  is the dimension of  $\mathbf{x}$ .

This update rule uses the step size  $\tau$  as a crude measure of proximity to the central path. In particular,  $\frac{2n}{\eta}$  is the value of  $t$  for which the associated central point has the same duality gap as the current point.

### 3.3.3 Termination Conditions

We considered multiple termination conditions during experimentation and although any one of these can be used, there is an additional cost associated with computing certain conditions and each condition also has to be calibrated appropriately so that convergence occurs even in the event of an inflection plateau.

#### Barrier Method

The barrier method uses a relative simple condition of  $\frac{n}{t} < \epsilon$  where  $n$  is equal to the number of inequality conditions (in this case, also the dimension of  $\mathbf{x}$ ). Note that  $\frac{n}{t}$  is the duality gap and hence can be used as a termination condition.

#### CIPM

As a stopping criterion, the method uses the duality gap divided by the dual objective value. By weak duality, the ratio is an upper bound on the relative suboptimality where  $p^*$  is the optimal value of the  $l_1$ -regularized LSP and  $f(\mathbf{x})$  is the primal objective computed with the point  $\mathbf{x}$ .

$$\frac{|f(\mathbf{x}) - p^*|}{|p^*|} \leq \frac{\eta}{G(\nu)}$$

Therefore we can use the duality term as the bound for termination:

$$\frac{\eta}{G(\nu)} < \epsilon$$

We can compute the dual feasible point  $\nu$ , dual objective  $G(\nu)$  and duality gap  $\eta$  as defined in ([Kim et al., 2007](#)).

$$\nu = 2\tau(\mathbf{A}\mathbf{x} - \mathbf{b})$$

The step size  $\tau$  computed from the line search is then updated to the following:

$$\tau = \min(\lambda / (2\|\mathbf{A}^T(\mathbf{A}\mathbf{x} - \mathbf{b})\|))$$

We can now compute the dual objective  $G(\nu)$  and the duality gap  $\eta$ :

$$G(\nu) = -\frac{1}{4}\nu^T\nu - \nu^T\mathbf{b}$$

$$\eta = \|\mathbf{A}\mathbf{x} - \mathbf{b}\|_2^2 + \lambda\|\mathbf{x}\|_1 - G(\nu)$$


#### Other

We also experimented with other termination conditions to check if these worked better.

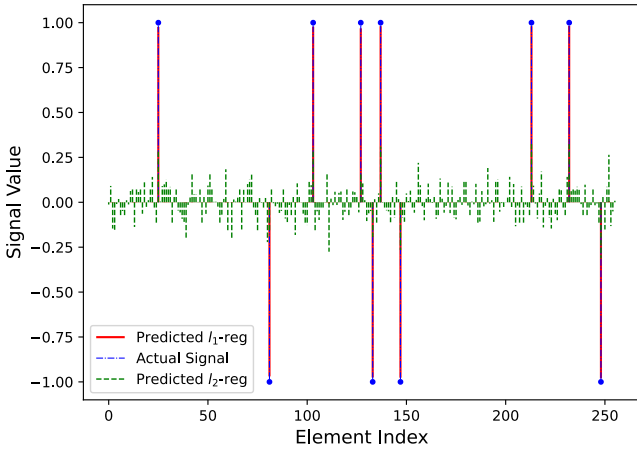
Similar to the derivation of CIPM, we can use a termination condition more commonly used for conjugate gradient methods which uses the gradients instead of the objective values.

$$\frac{|\nabla^2 f(\mathbf{x})\Delta\mathbf{x} + \nabla f(\mathbf{x})|}{|\nabla f(\mathbf{x})|} < \epsilon$$

We could also use the simplest condition usually used in gradient descent we could check if the magnitude of the gradient is less than some tolerance level which must be determined by experimentation.



(a) Barrier method (152 iterations)



(b) CIPM (573 iterations)

Figure 3.  $l_1$  signal reconstruction using NIPMs (CIPM and barrier method) with  $l_2$  analytical solution

$$|g| < \epsilon$$

Note that both of these failed to provide any improvements and in fact performed worse. The LASSO problem is not particularly suited for gradient based approaches and we find that the optimal solution is amidst large gradients.

### 3.3.4 Results

When we attempt to reconstruct the signal, we find that both methods find the optimal solution however it is important to note that the barrier method is able to find the optimum in fewer iterations with a lower computational cost than the CIPM defined in Kim et al. (2007) (Fig 3). It is unclear what benefits the more expensive  $t$ -update rule and the termination condition provides. The CIPM termination condition seemingly does not provide guaranteed convergence unlike the barrier method. Moreover the proposed truncated version of the algorithm actually uses inexact search directions for which the convergence is not proved. The PCG if implemented however would likely provide a performance improvement over computing the Hessian directly, especially for dimensionally larger signals.

### 3.4 Minimum Energy Comparison

To evaluate the performance of the solution for the  $l_1$ -regularised problem, we compare the result to the minimum energy reconstruction for the solution that is closest to the origin (ie. the  $l_2$ -regularised problem). This is also known as **ridge regression**.

$$\min_x \|Ax - b\|_2^2 + \lambda \|x\|_2^2$$

The analytical solution for this problem can be similarly derived as done for the least-squares problem (Eq 3).

$$\begin{aligned} f_{l_2}(x) &= \|Ax - b\|_2^2 + \lambda \|x\|_2^2 \\ &= (Ax - b)^T (Ax - b) + \lambda x^T x \\ &= x^T A^T A x - 2x^T A^T b + b^T b + \lambda x^T x \end{aligned}$$

$$\begin{aligned} \frac{d}{dx} f_{l_2}(x) &= 2A^T A x - 2A^T b + 2\lambda I x = 0 \\ (A^T A + \lambda I)x &= A^T b \\ \therefore x_{l_2} &= (A^T A + \lambda I)^{-1} A^T b \end{aligned} \quad (3)$$

Note that  $A^T A + \lambda I$  is always invertible for  $\lambda > 0$  and hence this solution is directly computable.

Also note that an equivalent solution  $x_{l_2}$  via the use of Lagrange multipliers gives  $x_{l_2} = A^T (A A^T)^{-1} b$  (Stankovic et al. (2019) Pg 21).

Solution Objective	$x_{l_1}$ BM	$x_{l_1}$ CIPM	$x_{l_2}$	$x_{opt}$
$f_{l_1}(x)$	0.06722	0.06805	0.13551	0.06767
$f_{l_2}(x)$	0.06624	0.06499	0.01797	0.06767

Table 2. Minimised values of the  $l_1$  and  $l_2$  regularised problem for the different solutions

From Table 2, we see that the  $l_2$  solution is far from the optimal solution of the  $l_1$  problem. We notice that the true  $l_1$  solution  $x_{opt}$  has the same function value for both objective functions which is optimal for the  $l_1$  problem but sub-optimal for the  $l_2$  problem. Therefore it is clear that the global minimum of the  $l_2$ -regularised problem is not the same as the global minimum of the  $l_1$ -regularised problem as seen in Figure 3 and the corresponding Table 2.

This is because the standard ridge regression, based on the  $l_2$ -norm, minimizes the energy of solution  $x$  and not its sparsity and hence this is not useful for reconstructing the desired signal as the signal has very few non-zero components. The  $l_1$  norm regularisation penalty puts the most weight on small residuals and the least weight on large values hence encouraging sparsity.

Table 2 shows that  $x_{l_1}$  BM is marginally better than  $x_{l_1}$  CIPM. Including the lower computational cost for BM, BM performs much better for sparse reconstruction in this specific case.



## 4 Discussion

The Newton method we used to perform signal reconstruction is in fact very inefficient and alternative methods such as iterative soft thresholding algorithm (ISTA) are computationally much cheaper. ISTA is specifically designed for reconstruction of sparse signals for the LASSO problem.

We create an implementation of ISTA and compare the results with the 2 NIPMs we investigated.

$$\text{soft}(x) = \text{sgn}(x) \max(0, |x| - \frac{\lambda}{2\alpha})$$

$$\mathbf{x}_{k+1} = \text{soft} \circ \left( \mathbf{x}_k + \frac{1}{\alpha} \mathbf{A}^T (\mathbf{b} - \mathbf{A} \mathbf{x}_k) \right)$$

The parameter  $\alpha$  is usually defined as  $2 \max(\text{eig}(\mathbf{A}^T \mathbf{A}))$ . The soft function is the soft-thresholding rule and the iteration step can be used as many times as required to converge to a solution. Hence this method is known as the iterative soft-thresholding algorithm (ISTA). Note that this is just one of possible solutions of the minimization problem with the  $l_1$ -norm (Stankovic et al. (2019) Pg 20).




Figure 4. Investigating  $\lambda$  for reconstruction

At the point of minimum reconstruction error  $\lambda \approx 0.02$  (Fig 4), we find that we require very few iterations for the algorithm to converge to the optimal solution in Figure 5. This is a good improvement over NIPMs both computationally as well as for accuracy.

Given by the decrease in error in Figure 5, we also see that the solution found is indeed optimal for our choice of  $\lambda$ . However this still requires more iterations than the barrier method for optimal convergence but it is important to note that each iteration is much more efficient than the barrier method. It is also important to note that we do not need to pick the  $\lambda$  that minimises the reconstruction error to find the optimal solution (as this requires knowledge of the optimal solution in the first place) but  $\lambda$  still needs to be chosen appropriately ( $0.003 \leq \lambda \leq 0.3$ ).

## References

Lieven Vandenberghe. Central Path, 2013. URL <http://www.seas.ucla.edu/~vandenbe/ee236a/lectures/cpath.pdf>.

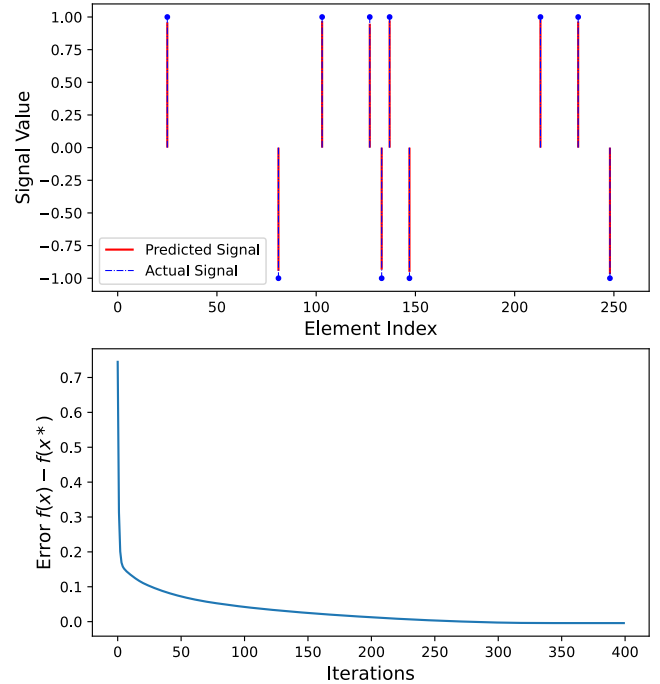


Figure 5. Reconstructed signal using ISTA for  $\lambda \approx 0.02$  (400 iterations)

Stephen Boyd and Lieven Vandenberghe. *Convex optimization*. Cambridge university press, 2004.

Geoff Gordon and Ryan Tibshirani. Gradient Descent Revisited, 2012. URL <https://www.cs.cmu.edu/~ggordon/10725-F12/slides/05-gd-revisited.pdf>.

Dmitriy Leykekhman. Regularized Linear Least Squares, 2018. URL [https://www2.math.uconn.edu/~leykekhman/courses/MATH3795/Lectures/Lecture\\_10\\_Linear\\_least\\_squares\\_reg.pdf](https://www2.math.uconn.edu/~leykekhman/courses/MATH3795/Lectures/Lecture_10_Linear_least_squares_reg.pdf).

Seung-Jean Kim, K. Koh, M. Lustig, Stephen Boyd, and Dimitry Gorinevsky. An interior-point method for large-scale  $l_1$ -regularized least squares. *IEEE Journal of Selected Topics in Signal Processing*, 1(4):606–617, 2007. doi: 10.1109/JSTSP.2007.910971.

Xiang Huang, Kuan He, Seunghwan Yoo, Oliver Cossairt, Aggelos Katsaggelos, Nicola Ferrier, and Mark Hereld. An interior point method for nonnegative sparse signal reconstruction. In *2018 25th IEEE International Conference on Image Processing (ICIP)*, pages 1193–1197, 2018. doi: 10.1109/ICIP.2018.8451710.

Danny Bickson and Danny Dolev. Distributed sensor selection using a truncated newton method, 2010.

Ljubisa Stankovic, Ervin Sejdic, Srdjan Stankovic, Milos Dakovic, and Irena Orovic. A tutorial on sparse signal reconstruction and its applications in signal processing. *Circuits, Systems, and Signal Processing*, 38, 03 2019. doi: 10.1007/s00034-018-0909-2.

## A Appendix

All Python files are available on [Google Drive](#).

```
1 import scipy.io as sio
2 import scipy.optimize
3 import numpy as np
4 import time
5
6 data_folder = "./data-q1/"
7
8 for iter in range(1, 6):
9     A_mat = "A"+str(iter)
10    b_arr = "b"+str(iter)
11    A1 = sio.loadmat(data_folder+A_mat+".mat")[A_mat]
12    b1 = sio.loadmat(data_folder+b_arr+".mat")[b_arr].squeeze()
13
14    M, N = A1.shape
15    #print(A1.shape)
16
17    I = np.eye(M)
18    One = np.ones((M,M))
19
20    c = np.concatenate((np.zeros(N), np.ones(M)))
21    A_ub_1 = np.vstack(( np.hstack((A1, -I)), np.hstack((-A1, -I)) )) # L_1
22    A_ub_inf = np.vstack(( np.hstack((A1, -One)), np.hstack((-A1, -One)) )) # L_inf
23    b_ub = np.concatenate((b1, -b1))
24
25    ###
26
27    # L_1 minimisation
28    start_1 = time.perf_counter()
29    res_1 = scipy.optimize.linprog(c, A_ub=A_ub_1, b_ub=b_ub)
30    end_1 = time.perf_counter() - start_1
31
32    # L_inf minimisation
33    start_inf = time.perf_counter()
34    res_inf = scipy.optimize.linprog(c, A_ub=A_ub_inf, b_ub=b_ub)
35    end_inf = time.perf_counter() - start_inf
36
37    # L_2 minimisation
38    start_2 = time.perf_counter()
39    x, sse, rank, s = np.linalg.lstsq(A1, b1)
40    end_2 = time.perf_counter() - start_2
41
42    ###
43
44    print(f"A{iter} b{iter}")
45    # Min values
46    print(f"Min Values")
47    print(round(res_1.fun, 2), round(np.sqrt(sse[0]), 2),
48          round(res_inf.fun, 2))
49
50    # Runtime
51    print(f"Runtime")
52    print(round(end_1*1000, 1), round(end_2*1000, 1),
53          round(end_inf*1000, 1))
54
55    print("#####")
```

Listing 1. Norm Table - Generate Table 1

```
1 import scipy.io as sio
2 import scipy.optimize
3 import scipy.stats
4 import numpy as np
5 import matplotlib.pyplot as plt
6
7 data_folder = "./data-q1/"
8 A_mat = "A5"
9 b_arr = "b5"
10 A1 = sio.loadmat(data_folder+A_mat+".mat")[A_mat]
11 b1 = sio.loadmat(data_folder+b_arr+".mat")[b_arr].squeeze()
12
13 M, N = A1.shape
14 #print(A1.shape)
15
16 I = np.eye(M)
17 One = np.ones((M,M))
18
19 c = np.concatenate((np.zeros(N), np.ones(M)))
20 A_ub_1 = np.vstack(( np.hstack((A1, -I)), np.hstack((-A1, -I)) )) # L_1
21 A_ub_inf = np.vstack(( np.hstack((A1, -One)), np.hstack((-A1, -One)) )) # L_inf
22 b_ub = np.concatenate((b1, -b1))
23
24 #print(A_ub_1.shape) # 2m x (n+m)
25 #print(b_ub.shape) # 2m
26 #print(c.shape) # n+m
27
28 res_1 = scipy.optimize.linprog(c, A_ub=A_ub_1, b_ub=b_ub)
29 # L_1 minimisation
30
31 res_inf = scipy.optimize.linprog(c, A_ub=A_ub_inf, b_ub=b_ub)
32 # L_inf minimisation
33
34 x, sse, rank, s = np.linalg.lstsq(A1, b1) # L_2
35 minimisation
36
37 resid_1 = np.dot(A1, res_1.x[:N]) - b1
38 resid_inf = np.dot(A1, res_inf.x[:N]) - b1
39 resid_2 = np.dot(A1, x) - b1
40
41
42 xmin, xmax, xcount = -2, 2, 1000
43 lnspec = np.linspace(xmin, xmax, xcount)
44
45 #####
46
47 # Laplace (1-Norm)
48 ag, bg = scipy.stats.laplace.fit(resid_1)
49 pdf_laplace = scipy.stats.laplace.pdf(lnspec, ag, bg)
50
51 # Gaussian (2-Norm)
52 m, s = scipy.stats.norm.fit(resid_2)
53 pdf_g = scipy.stats.norm.pdf(lnspec, m, s)
54
55 #####
56
57 num_bins = 40
58 area_scale_1 = (max(resid_1)-min(resid_1))/num_bins * len(resid_1)
59 area_scale_2 = (max(resid_2)-min(resid_2))/num_bins * len(resid_2)
60 #area_scale_inf = (max(resid_inf)-min(resid_inf))/num_bins * len(resid_inf)
61
62 fig, axs = plt.subplots(3, 1, sharex=True)
63
64 axs[0].hist(resid_1, bins=num_bins, density=False, label=r"$L_1$", edgecolor='b', color='w')
65 axs[0].plot(lnspec, area_scale_1*pdf_laplace, label="Laplace", color='b')
66
67 axs[1].hist(resid_2, bins=num_bins, density=False, label=r"$L_2$", edgecolor='b', color='w')
68 axs[1].plot(lnspec, area_scale_2*pdf_g, label="Gaussian", color='b')
69
70 axs[2].hist(resid_inf, bins=num_bins, density=False, label=r"$L_{\infty}$", edgecolor='b', color='w')
71
72 axs[0].set_ylabel("Frequency")
73 axs[1].set_ylabel("Frequency")
74 axs[2].set_ylabel("Frequency")
75
76 axs[0].set_xlabel("Residual")
77
78 plt.show()
```

Listing 2. Norm Residuals - Generate Figure 1



```

1 import scipy.io as sio
2 import scipy.optimize
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 data_folder = "./data-q1/"
7 A.mat, b_arr = "A3", "b3"
8 A1 = sio.loadmat(data_folder+A.mat+".mat")[A.mat]
9 b1 = sio.loadmat(data_folder+b_arr+".mat")[b_arr].squeeze()
10 M, N = A1.shape
11
12 I = np.eye(M)
13 One = np.ones((M,M))
14
15 c = np.concatenate((np.zeros(N), np.ones(M)))
16 A = np.vstack(( np.hstack((-A1, -I)), np.hstack((-A1, -I)) ))
17 b = np.concatenate((b1, -b1))
18
19 def f(x):
20     global A, b, c
21     return np.dot(c, x) - np.sum(np.log(np.abs(b - A@x)))
22
23 def grad(x):
24     global A, b, c
25     d = 1/(b - A@x)
26     #print("Condition Number", np.linalg.cond(A.T @ (np.diag
27         (d)**2 @ A) ))
28     return c + A.T @ d
29
30 # Parameters.
31 alpha = 0.3
32 beta = 0.8
33 t = 1
34 np.random.seed(0)
35 #x = np.zeros(N+M)
36 x = np.ones(N+M)*0.1
37 iter = 1
38
39 func_vals = []
40 while iter < 5e4:
41     func_vals.append(f(x))
42
43     g_vec = grad(x)
44     # set descent direction to negative normalised gradient
45     g = -g_vec/np.linalg.norm(g_vec)
46
47     if iter % 1000 == 0:
48         print(f'[{iter}] \t Grad: {np.linalg.norm(g_vec)} \t
49             Step-size: {t}')
50
51     # Backtracking linesearch implementation
52     while f(x + t*g) > f(x) + (alpha * t * np.dot(g_vec, g))
53         :
54         t *= beta
55
56     # Exact linesearch
57     #res = scipy.optimize.line_search(f, grad, x, g)
58     #t = res[0]
59
60     x += t * g
61     iter += 1
62
63 log_err = np.abs(np.array(func_vals) - func_vals[-1])
64 fig, axs = plt.subplots()
65 axs.plot(log_err)
66 #axs.plot(log_err[:40000])
67 #axs.set_ylim(1e2, 1e4)
68
69 axs.set_yscale('log')
70 axs.set_ylabel(r"$|f(x_k) - p^*|$", fontsize="13")
71 axs.set_xlabel("Iterations", fontsize="13")
72 plt.show()

```

Listing 3. Backtracking Linesearch for  $l_1$  Norm Approximation - Generate Fig 2

```

1 import scipy.io as sio
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 data_folder = "./data-q3/"
6 A.mat, x_vec = "A", "x0"
7 A = sio.loadmat(data_folder+A.mat+".mat")[A.mat]
8 xopt = sio.loadmat(data_folder+x_vec+".mat")[x_vec].squeeze()
9
10 M, N = A.shape
11 b = np.dot(A, xopt)
12 lam_max = np.linalg.norm(2*A.T@b, np.inf)
13 lam = 0.01 * lam_max
14
15 def objective_l2(x):
16     global A, b, lam
17     return np.linalg.norm(A@x - b)**2 + lam*np.linalg.norm(x, 2)**2
18
19 def objective(x):
20     global A, b, lam
21     return np.linalg.norm(A@x-b)**2+lam*np.linalg.norm(x,1)
22
23 def f(x, u, t):
24     global A, b, lam
25     return t*(np.linalg.norm(A@x - b)**2 + lam*np.sum(u)) -
26         np.sum(np.log(np.abs(u**2-x**2)))
27
28 def grad_x(x,u,t):
29     global A, b, lam
30     return 2* (t * A.T @ (A@x-b) + x / (u**2 - x**2))
31
32 def grad_u(x,u,t):
33     global A, b, lam
34     return t*lam*np.ones(len(u)) - 2 * u/(u**2 - x**2)
35
36 def hessian(x,u,t):
37     global A, b, lam
38     den = (u**2-x**2)
39     D1 = np.diag( 2* ((u**2+x**2) / den) / den )
40     D2 = np.diag( ((-4*x*u) / den) / den )
41     H = np.vstack(( np.hstack((2*t*A.T@A + D1, D2)), np.
42         hstack((D2, D1)) ))
43     return H
44
45 t = 1/lam
46 mu, epsilon = 1.1, 1e-6
47 alpha, beta, tau = 0.25, 0.4, 1
48 x, u = np.zeros(N), np.ones(N)
49 g_vec_x, g_vec_u = grad_x(x, u, t), grad_u(x, u, t)
50
51 iter = 1
52 while iter <= 200:
53     g_vec_x, g_vec_u = grad_x(x, u, t), grad_u(x, u, t)
54     g_vec = np.hstack(( g_vec_x, g_vec_u ))
55     h = hessian(x, u, t)
56
57     # Compute Newton search direction
58     g = np.linalg.solve(h, -g_vec)
59     g /= np.linalg.norm(g)
60     g_x, g_u = g[:N], g[N:]
61
62     # Backtracking linesearch implementation
63     while f(x + tau*g_x, u+tau*g_u, t) > f(x, u, t) + (alpha
64         * tau * (np.dot(g_vec_x, g_x) + np.dot(g_vec_u, g_u))
65         ):
66         tau *= beta
67
68     x += tau * g_x
69     u += tau * g_u
70     print(f"Iter {iter} \t Step-size {tau} \t t {t}")
71     print("N/t", N/t)
72
73     if N/t < epsilon:
74         break
75
76     t *= mu
77     iter += 1
78
79 xnew = x
80
81 # Both solutions are equivalent
82 x_l2 = A.T@(np.linalg.inv(A@A.T@b)
83     #x_l2 = np.linalg.inv(A.T@A + lam*np.eye(N))@(A.T@b)
84
85 fig, axs = plt.subplots(figsize=(7,5))
86 axs.vlines(range(N), ymin=np.zeros(N), ymax=xnew, colors=["r"]
87     *N, label="Predicted "+r"$1.1$"+"-reg", lw=1.5)
88 axs.vlines(range(N), ymin=np.zeros(N), ymax=xopt, colors=["b"]
89     *N, label="Actual Signal", linestyle="-. ", lw=0.75)
90
91 nz_inds = np.argwhere(xopt != 0)
92 axs.scatter(nz_inds, xopt[nz_inds], c="b", marker=".")
93
94 # L2 Prediction
95 axs.vlines(range(N), ymin=np.zeros(N), ymax=x_l2, colors=["g"]
96     *N, label="Predicted "+r"$1.2$"+"-reg", linestyle="-. ", lw=1)
97
98
99 axs.legend(loc="lower left")
100 axs.set_xlabel("Element Index", fontsize=13)
101 axs.set_ylabel("Signal Value", fontsize=13)
102 plt.show()

```

Listing 4. BM - Generate Fig 3a

```

1 import scipy.io as sio
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 data_folder = "./data-q3/"
6 A_mat, x_vec = "A", "x0"
7 A = sio.loadmat(data_folder+A_mat+".mat")[A_mat]
8 xopt = sio.loadmat(data_folder+x_vec+".mat")[x_vec].squeeze()
9 M, N = A.shape
10 b = np.dot(A, xopt)
11 lam_max = np.linalg.norm(2*A.T@b, np.inf)
12 lam = 0.01 * lam_max
13
14 def objective_l2(x):
15     global A, b, lam
16     return np.linalg.norm(A@x - b)**2 + lam*np.linalg.norm(x, 2)**2
17
18 def objective(x):
19     global A, b, lam
20     return np.linalg.norm(A@x-b)**2+lam*np.linalg.norm(x, 1)
21
22 def dual_objective_G(v):
23     global A, b, lam
24     return -(1/4) * v@v - v@b
25
26 def f(x, u, t):
27     global A, b, lam
28     return t*(np.linalg.norm(A@x - b)**2 + lam*np.sum(u)) - np.sum(np.log(np.abs(u**2-x**2)))
29
30 def grad_x(x,u,t):
31     global A, b, lam
32     return 2* (t* A.T @ (A@x-b) + x / (u**2 - x**2))
33
34 def grad_u(x,u,t):
35     global A, b, lam
36     return t*lam*np.ones(len(u)) - 2 * u/(u**2 - x**2)
37
38 def hessian(x,u,t):
39     global A, b, lam
40     den = (u**2-x**2)
41     D1 = np.diag( 2* ((u**2+x**2) / den) / den )
42     D2 = np.diag( ((-4*x*u) / den) / den )
43     H = np.vstack(( np.hstack((2*t*A.T@A + D1, D2)), np.hstack((D2, D1)) ))
44     return H
45
46 t, mu, epsilon = 1/lam, 1.1, 0.2
47 alpha, beta, tau, tau_min = 0.01, 0.4, 1, 0.5
48 x, u = np.zeros(N), np.ones(N)
49 g_vec_x, g_vec_u = grad_x(x, u, t), grad_u(x, u, t)
50
51 iter = 1
52 while iter <= 1000:
53     g_vec_x, g_vec_u = grad_x(x, u, t), grad_u(x, u, t)
54     g_vec = np.hstack(( g_vec_x, g_vec_u ))
55     h = hessian(x, u, t)
56     # Compute Newton search direction
57     g = np.linalg.solve(h, -g_vec)
58     g /= np.linalg.norm(g)
59     g_x, g_u = g[:N], g[N:]
60
61     # Backtracking linesearch implementation
62     while f(x+tau*g_x,u+tau*g_u,t) > f(x,u,t) + (alpha*tau*(np.dot(g_vec_x, g_x)+np.dot(g_vec_u, g_u))):
63         tau *= beta
64
65     x += tau * g_x
66     u += tau * g_u
67     v = 2*tau*(A@x - b) # dual feasible point v
68     tau = min(lam/np.abs(2*A.T@(A@x - b)))
69     G_v = dual_objective_G(v)
70     eta = np.abs(objective(x) - G_v) # duality gap eta
71
72     print(f"Iter {iter} \t Step-size {tau} \t t {t}")
73     print("Term Cond", eta/G_v)
74     if eta/G_v < epsilon:
75         break
76     if tau >= tau_min:
77         t = max( mu * 2*min(N/eta, t), t )
78     iter += 1
79
80 xnew = x
81 # Both solutions are equivalent
82 x_l2 = A.T@(np.linalg.inv(A@A.T)@b)
83 #x_l2 = np.linalg.inv(A.T@A + lam*np.eye(N))@(A.T@b)
84
85 fig, axs = plt.subplots(figsize=(7,5))
86 zs = np.zeros(N)
87 axs.vlines(range(N), ymin=zs, ymax=xnew, colors=["r"]*N, label="Predicted "+r"$\hat{x}_l$"+"-reg", lw=1.5)
88 axs.vlines(range(N), ymin=zs, ymax=xopt, colors=["b"]*N, label="Actual Signal", linestyle='-', lw=0.75)
89
90 nz_inds = np.argwhere(xopt != 0)
91 axs.scatter(nz_inds, xopt[nz_inds], c="b", marker=".")
92
93 axs.vlines(range(N), ymin=zs, ymax=x_l2, colors=["g"]*N, label="Predicted "+r"$\hat{x}_l$"+"-reg", linestyle='--', lw=1)
94
95 axs.legend(loc="lower left")
96 axs.set_xlabel("Element Index", fontsize=13)
97 axs.set_ylabel("Signal Value", fontsize=13)
98 plt.show()

```

Listing 5. CIPM - Generate Fig 3b

```

1 import scipy.io as sio
2 import scipy.optimize
3 import scipy.stats
4 import numpy as np
5 import matplotlib.pyplot as plt
6
7 data_folder = "./data-q3/"
8 A_mat, x_vec = "A", "x0"
9 A = sio.loadmat(data_folder+A_mat+".mat")[A_mat]
10 xopt = sio.loadmat(data_folder+x_vec+".mat")[x_vec].squeeze()
11 M, N = A.shape
12 b = np.dot(A, xopt)
13 lam_max = np.linalg.norm(2*A.T@b, np.inf)
14
15 # primal objective value of x
16 def objective(x):
17     global A, b, lam
18     return np.linalg.norm(A@x - b)**2 + lam*np.linalg.norm(x, 1)
19
20 def soft(y, lambd):
21     c = np.maximum(np.zeros(N), np.abs(y) - lambd)
22     return np.sign(y) * c
23
24 np.random.seed(0)
25 # Note alpha is required to be greater than lam_max
26 alpha = 2*np.max(np.linalg.eigvals(A.T@A)).real # this is commonly used
27 #print("alpha", alpha, "\t lam_max", lam_max)
28
29 x = (np.random.rand(N)*2-1)*0.01
30
31 xs = []
32 xopts = []
33 factors = np.array(range(-20, 4))
34 lams = (10**factors/4) * lam_max
35
36 for lam in lams:
37     iter = 1
38     while iter <= 200:
39         y = 1/alpha * A.T@(b-A@x) + x
40         x = soft(y, lam/(2*alpha))
41         #print(lam)
42         iter += 1
43     xs.append(x)
44     xopts.append(xopt)
45
46 fig, axs = plt.subplots()
47 axs.plot(lams, np.linalg.norm(np.array(xs)-np.array(xopts), axis=1))
48 #print(lams[np.argmin(np.linalg.norm(np.array(xs)-np.array(xopts), axis=1))])
49
50 axs.set_xscale('log')
51 axs.set_xlabel(r"$\lambda$", fontsize=13)
52 axs.set_ylabel("Reconstruction Error "+r"$||x-x^*||_2$", fontsize=13)
53 plt.show()

```

Listing 6. ISTA  $\lambda$ -Variation - Generate Fig 4

```

1 import scipy.io as sio
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 data_folder = "./data-q3/"
6 A_mat, x_vec = "A", "x0"
7 A = sio.loadmat(data_folder+A_mat+".mat")[A_mat]
8 xopt = sio.loadmat(data_folder+x_vec+".mat")[x_vec].squeeze()
9 M, N = A.shape
10 b = np.dot(A, xopt)
11 lam_max = np.linalg.norm(2*A.T@b, np.inf)
12 lam = 0.03 * lam_max
13
14 # primal objective value of x
15 def objective(x):
16     global A, b, lam
17     return np.linalg.norm(A@x - b)**2 + lam*np.linalg.norm(x, 1)
18
19 def soft(y, lamb):
20     c = np.maximum(np.zeros(N), np.abs(y) - lamb)
21     return np.sign(y) * c
22
23 np.random.seed(0)
24 # Note alpha is required to be greater than lam_max
25 alpha = 2*np.max(np.linalg.eigvals(A.T@A)).real # this is commonly used
26 print("lam", lam, "\t alpha", alpha, "\t lam_max", lam_max)
27
28 x = np.zeros(N)
29 #x = (np.random.rand(N)*2-1)*0.01
30
31 func_vals = []
32 iter = 1
33 while iter <= 400:
34     y = 1/alpha * A.T@(b-A@x) + x
35     x = soft(y, lam/(2*alpha))
36     func_vals.append(objective(x))
37     #print(f"Iter {iter}")
38     iter += 1
39
40 xnew = x
41 print(objective(xnew), objective(xopt))
42
43 fig, axs = plt.subplots(2,1, figsize=(7,8))
44 # Plot components as vertical signal lines
45 axs[0].vlines(range(N), ymin=np.zeros(N), ymax=xnew, colors=["r"]*N, label="Predicted Signal", lw=1.5)
46 axs[0].vlines(range(N), ymin=np.zeros(N), ymax=xopt, colors=["b"]*N, label="Actual Signal", linestyle='-.', lw=0.75)
47
48 # For aesthetics and visibility of actual signal
49 non_zero_indices = np.argwhere(xopt != 0)
50 axs[0].scatter(non_zero_indices, xopt[non_zero_indices], c="b", marker=".")
51
52 axs[1].plot(np.array(func_vals)-objective(xopt))
53
54 axs[0].legend(loc="lower left")
55 axs[0].set_xlabel("Element Index", fontsize=13)
56 axs[0].set_ylabel("Signal Value", fontsize=13)
57
58 axs[1].set_xlabel("Iterations", fontsize=13)
59 axs[1].set_ylabel("Error "+r"$f(x) - f(x^*)$", fontsize=13)
60
61 plt.show()

```

Listing 7. ISTA Method - Generate Fig 5